

Caregiver reconstruction of children's errors: the preservation of complexity in language (improve title)

Anonymous CogSci submission

Abstract

Why do languages change? One possibility is they evolve due to two competing pressures: one, for the language to be easily transmitted to new generations—and hence simple—and another, for the language to be a useful, descriptive form of communication—and hence more complex. However, few studies have explored these pressures in the most frequent language learners: children. Conventional iterated learning studies focus on the transmission of a novel language from one adult to the next. However, this ignores one of the important features of language learning, namely, receiving feedback. In this study, 120 adults on Mechanical Turk participated in a conventional iterated learning task. Their results were replicated with a larger sample of 480 adults. Participants' performances indicate that complexity decreases in the language over generations of transmission. 960 adults participated in a third experiment, where half of the participants represented parent or caregivers, who added a process of editing into generational transmission (this all needs to be rewritten and re-worded). Results show that adding this error-correcting participant allows a greater level of complexity to be retained in the language compared with having a sole language learner/transmitter.

Keywords: communication; language acquisition; language evolution; iterated learning

Introduction

How do you ask a group of people where they are going in Spanish? In Spain, the answer depends on the group: you might ask “Donde van ustedes?” of a group of work colleagues, but to address your friends, you use the informal “Donde váis vosotros?” instead. In Mexican Spanish, this distinction has disappeared, and the “ustedes” form is used exclusively. Why did Spanish change in this way, simplifying and shedding the formal second person plural? Why do languages change at all? One working theory is that languages evolve to adapt to two dynamic competing pressures: one, to be easily transmitted and learned (and hence simple), and another, to be a descriptive and effective system for communication (and hence complex) (Lupyan & Dale, 2010).

When children are learning language, they often make simplification errors (Bowerman, 1982). If a child is asking for some milk, they may ask for a “baba” (bottle). Thus, the child has coined a new word, which her parents understand. This simplification works well for the child, until she starts to learn her animal words, and calls a sheep “baba”. The child language-learner has shown the effects of the simplicity pressure in language: if she calls both bottles and sheep “baba”, she has to learn one less word. Yet, this simplicity bias causes

problems when she asks her parent for a “baba” – does she want the sheep toy or her bottle? Her parent will then attempt to resolve this ambiguity, by introducing descriptiveness into the child's vocabulary (“Do you want the sheep or the baba (bottle)?”). Due to the parent's intervention, this child will not grow up to think that “sheep” and “bottles” can be named with the same label. [this paragraph is really long and long-winded, definitely simplify but high-level, is this on the right track? Do I want to include this extended example? Should I start the intro with this and scrap the Spanish example??]

Children are often the actors who drive language evolution (Senghas, 2003), yet they differ from adults in their cognitive capabilities, namely, memory systems (Kempe, Gauvrit, & Forsyth, 2015), interests and early vocabularies, and conversation partners. Therefore, though children are skilled language learners, their developing cognitive systems prevent aspects of language that are difficult to learn and remember from being passed on—pushing languages towards simplicity (Hudson Kam & Newport, 2005; Senghas, 2003). But, languages that become too simple can lose the ability to be effective for communication (Kirby, Griffiths, & Smith, 2014). What enables languages to retain their communicative utility in the face of these learnability pressures?

The following study tests a novel hypothesis for the maintenance of structure in language: Communicative inference by caregivers. [is “communicative inference” clear from the rest of this paragraph? Does it need to be better-defined? How so?] Children's language learning is greatly influenced by those around them—especially the adults they talk to most. These caregivers determine the majority of their child's language input, and are responsible for seeing that their children develop effective and useful systems of communication. Even the youngest children are not passive learners of language — they are active participants, engaging in conversations with their parents. These adults are experts both in the language and in the children themselves, as they understand the child's intuitions, personality, and context. Caregivers play an important interpretive role in these interactions through their ability to understand the intended target of children's errorful productions (Chouinard & Clark, 2003). They may reformulate their child's language either through explicit or implicit correction. Adults can explicitly correct their children's errors in various ways (e.g., by interruptions or repeating the correct word/grammatical form) (Penner, 1987). Yet, children

primarily learn language through listening to others talk, instead of being explicitly instructed in a language (Romberg & Saffran, 2010). Thus, parent's modeling of accurate language constructions can have a powerful effect on reducing children's language errors: over time, children fix their own mistakes because they have learned the correct constructions from their caregivers (Hudson Kam & Newport, 2005). By way of this feedback, both implicit and explicit, children's simplification errors are corrected, and children are able to acquire adult-like speech. Eventually, when a child grows to be an adult, they will not transmit the errors they had as a child, but the correct forms of speech they learned from their caregivers. Thus, over the course of a lifetime, the child language learner grows to become a parent language teacher, correcting their own children's errors. These error reconstructions may be a mechanism by which more structure is retained in language over many lifetimes than children could sustain alone.

[okay transition?] Although language evolution is a process which is constantly shaping our languages, studying the pressures of simplicity and descriptiveness on language evolution is relatively difficult. This is partially due to the complexity of naturalistic language change (Ellis, 2008). There are also few instances of naturalistic language emergence in speakers without a prior language (Brentari & Coppola, 2013; Sandler, Meir, Padden, & Aronoff, 2005; Senghas, 2003). Therefore, researchers have developed methods for studying language emergence, evolution, and change within the laboratory. [this is a bad ending sentence]

Using iterated learning to study language change

One way to study language acquisition in the lab is to use the iterated learning paradigm. This paradigm was created to study the effects of simplicity and informativity on inter-generational language evolution (Kirby et al., 2014). In an iterated learning paradigm, one participant is trained on a randomly-generated language—for example, a set of words created by arbitrarily pairing syllables together. The participant is later asked to recall the language, and their responses are given as training input for the next subject, thus creating a transmission chain. This iterated process mimics the transmission of language across generations, with each participant unintentionally changing the language through their memory biases. However, the vast majority of iterated learning studies have used adult participants (Christiansen & Kirby, 2003; Kirby, Dowman, & Griffiths, 2007; Kirby et al., 2014; e.g., Smith & Wonnacott, 2010; structure & signals, 2014), and few studies have used children as research subjects (Kempe et al., 2015; Raviv & Arnon, 2018). A study by Kempe et al. (2015) compared child (age 5-8) and adult performances on an iterated learning task using a novel dot-pattern paradigm. Their results found that structure emerged faster in children than adults—that is, the children's patterns simplified much faster than the adult's, allowing them to be easily reproducible earlier in the transmission chain. This study provides evidence of the importance of looking at both children and adults in an iterated learning paradigm, as they have

different cognitive skills which affect their performance in language-learning tasks (Kempe et al., 2015). However, language evolution cannot be fully grasped using this paradigm with only separate adult or child learning chains, because language learning does not occur only within the same age group (horizontal transmission), or only across age groups (vertical transmission), but it occurs dynamically, in both directions. In a true language-acquisition environment, a child receives both language input and feedback from their caregiver and uses it to interact with their peers throughout life, eventually growing into a new teacher-caregiver.

The following study uses an iterated learning (diffusion chain) paradigm with an adaptation of Kempe et al. (2015)'s original non-linguistic task to investigate how the pressures of descriptiveness and transmissibility operate in adults when their responses are subject to error correction by a secondary participant. We hypothesize that these error-correctors (analogous to caregivers and teachers) are pivotal not only to an individual's successful language acquisition, but also to the evolution of a language as a whole. This is because those who correct mistakes and provide feedback are able to protect against the transmissibility (simplicity) bias, which is likely stronger in early language learners, by re-introducing and preserving complexity in language.

Experiments 1 and 2 attempt to replicate Kempe et al. (2015)'s findings that, using an artificial dot-pattern language, complexity decreases over generations and transmission accuracy increases. This finding reflects a stronger bias of language-learners towards the simplicity pressure, rather than descriptiveness pressure. Experiment 3 investigates the addition of a parent-like participant into the generational transmission process: instead of language-learners transmitting their reproductions directly to the next generation, we include a secondary participant. This participant acts as a caregiver, attempting to compensate or fix some of the errors introduced by the reproducers to better reproduce the initial language. Just as adults implicitly and explicitly attempt to teach their children to produce accurate forms of their native language(s), these secondary participants are attempting to help the reproducers better produce accurate forms of this artificial language. This edited form of language is then passed to the next reproducing participant, thus representing a child learner who has grown and implemented feedback and error-correction provided by their parent over the years. We thus hypothesize that the addition of this secondary caregiver-like participant is pivotal to re-introducing and maintaining a higher level of descriptiveness (complexity) in language. [is this too much detail? not enough? clear analogy?]

Experiment 1: Replicating Kempe et al. (2015)

In a baseline experiment, adults participated in a standard iterated learning study, using stimuli adapted from Kempe et al. (2015). Participants were told to reproduce patterns on grids, and each user's responses were used as training input for the subsequent participant.

Method

Participants Participants in Experiment 1 were 120 adults recruited on Amazon Mechanical Turk. These participants were divided into twenty diffusion chains, each of which had six generations. Each participant gave informed consent. The task was approximately eight minutes long, and subjects were compensated with \$0.50 for their participation.

Design and Procedure Participants in Experiment 1 were told that in this task, they would be re-creating patterns on a grid. After a consent screen, subjects first viewed a training trial with two 8x8 grids on the screen – a target grid, with 10 cells colored in, and a blank grid. They were told to make the blank grid match the target grid exactly, and were unable to progress until the grids were identical. Following this trial, participants were informed that they would see a target grid appear on the screen for 10 seconds, followed by a picture (a visual mask) displayed for 3 seconds. After the visual mask, participants viewed a blank 8x8 grid where they were given 60 seconds to re-create the target grid. Participants could click on any cell in the grid to change its color, and could also undo any color placed. A counter on the screen showed how many targets had been colored, and it varied dynamically with the participant’s clicks. After placing 10 targets, participants could click a button to advance to the next trial. After completing 3 Practice trials, which were identical for all participants in all chains, participants were informed that the study would begin.

Each participant then completed 6 Experiment trials. Participants in the first generation of each chain received the same initial grids for Experiment trials. These initial 8x8 grids were generated by randomly selecting 10 of the 64 possible cells to be filled. Participants in subsequent chains received as their targets the outputs produced by their parent in the chain.

Participants’ performance on the practice trials were used as an attention check to determine whether their data would be passed to the next participant. If the participant scored less than 75% accuracy on the last 2 practice trials, or if they failed to select 10 cells before time ran out, their outputs were not transmitted to the next generation. In Experiment 1, five participants failed to meet these criteria and were excluded from analysis.

All experiments were pre-registered on Open Science Framework and can be accessed at the following link: <https://osf.io/guzyf/>.

Results and Analysis

Our primary measures of interest were reproduction accuracy and pattern complexity. Reproduction accuracy served as a proxy for transmissibility – higher reproduction accuracies indicate that the “language” is easier to learn. Reproduction accuracy was computed as the proportion of targets (out of 10) which were placed in the same location on the target and input grids.

Complexity served as a proxy for descriptiveness. We

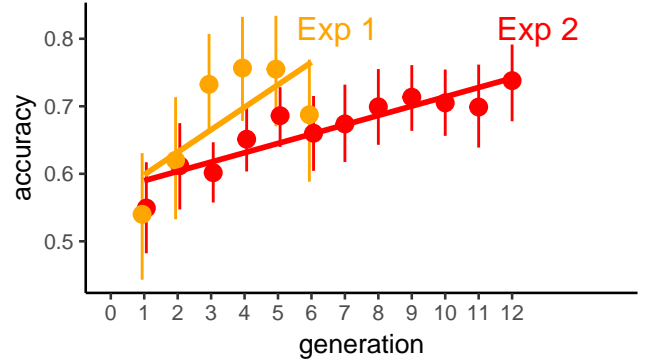


Figure 1: Experiments 1 and 2 show increases in accuracy over transmission generations. CHANGE POINT SIZE

followed Kempe et al. (2015) in using several measures of complexity: algorithmic complexity, chunking, and edge length. Algorithmic complexity is calculated using the Block Decomposition Method, a measure of Kolmogorov-Chaitin Complexity applied to 2-dimensional patterns (Zenil, Soler-Toscano, Dingle, & Louis, 2014). This measure computes the length of the shortest Turing machine program required to produce the observed pattern. The shorter the program, the simpler the pattern. Chunking is the number of groups of colored blocks which share an edge. The more groups of blocks, the easier the pattern is to transmit, and the lower the complexity is. Edge length is the total perimeter of the colored blocks. If all blocks were in one chunk, the edge length would be low, and the complexity of the pattern would likely be lower compared to if none of the chosen targets shared an edge. Implementation of these metrics was adapted from code provided by Gauvrit, Soler-Toscano, & Guida (2017).

Results

If iterated learning captures the hypothesized pressures of expressiveness and transmissibility, we predict that over generations reproduction accuracy should increase and complexity should decrease. We tested these predictions with mixed-effects logistic regressions, predicting accuracy and all three measures of complexity separately from fixed effects of generation and trial number, and random intercepts for participant, chain, and initial grid (e.g. $\text{accuracy} \sim \text{generation} + \text{trial} + (1|\text{subject}) + (1|\text{initial}) + (1|\text{chain})$).

Reproduction accuracy increased significantly over generations ($\beta = 0.025$, $t = 3.044$, $p = .003$). Figure 1 shows the results for accuracy. Complexity on all three measure decreased significantly over generations ($\beta_{BDM} = -3.219$, $t = -7.599$, $p < .001$; $\beta_{\text{chunking}} = -0.35$, $t = -12.671$, $p < .001$; $\beta_{\text{edge}} = -0.763$, $t = -12.256$, $p < .001$). Figure 2 shows the results for algorithmic complexity.

Experiment 2: Replication and extension of Experiment 1

Experiment 2, replicated the task from Experiment 1 with the addition of twice as many chains and generations. We repli-




Figure 2: Experiments 1 and 2 show decreases in algorithmic complexity over time. Grid patterns as produced by participants at generations 0 (initial), 5, and 11 are shown.

cated the task with a larger sample in order to approximate the shape of the algorithmic complexity curve. Particularly, we were interested in whether complexity asymptoted over generations.

Method

Participants

Participants in Experiment 2 were 519 adults recruited on Amazon Mechanical Turk. These participants were divided into forty diffusion chains, each of which had twelve generations. Each participant gave informed consent, and was compensated with \$0.50 for their participation.

Design and Procedure

The task in Experiment 2 was identical to Experiment 1. Participants were told to reproduce patterns on a grid, and their responses were passed to the next subject in the transmission chain.

Approximately 8% ($n=39$) of participants in Experiment 2 were excluded from analysis due to failure to meet accuracy requirements on the practice trials or failure to select the complete number of targets on one or more experimental trials. This resulted in a total of 480 participants included in the analysis.

Results

The results of this experiment replicated those found in Experiment 1. Reproduction accuracy increased significantly over generations ($\beta = 0.051$, $t = 4.378$, $p = < .001$). Figure 1 shows the results for accuracy.

Figure 2 shows the results for algorithmic complexity. Algorithmic complexity appeared to follow (W.C.??) an exponential function of the form $y = e^{-x} + b$. We therefore fit an exponential mixed-effects regression model predicting complexity from fixed effects of generation and trial number, and random intercepts for participant, chain, and initial grid (e.g. $\log(\text{complexity}) \sim \text{generation} + \text{trial} +$

$(1|\text{subject}) + (1|\text{initial}) + (1|\text{chain})$). Algorithmic complexity decreased and asymptoted over generations ($\beta_{BDM} = -0.039$, $t = -5.202$, $p = < .001$). Similar trends were also found with chunking and edge length, the alternate measures of complexity (CHECK THESE MODELS B/C THE LOG EXPONENTIAL DOESN'T WORK $\beta_{\text{chunking}} = -0.834$, $t = -14.47$, $p = < .001$; $\beta_{\text{edge}} = -1.601$, $t = -10.872$, $p = < .001$).

Experiment 3: Introducing an interlocutor

In order to add an element of feedback from a more experienced interlocutor to the iterated-learning process, we adapted the task from Experiments 1 and 2 to include a secondary, “editing” participant. This participant was analogous to a caregiver who protects their child from learning incorrect forms of language.

Method

Participants

Participants in Experiment 3 were 1031 adults recruited on Amazon Mechanical Turk. These participants were divided into forty diffusion chains, each of which had twelve generations. Each participant gave informed consent, and was compensated with \$0.50 for their participation.

Design and Procedure

In the third, dyad experiment, a primary participant was designated to be a “learner”, and completed the same task as in Experiment 1 and Experiment 2. They were told to reproduce patterns on a grid. A secondary participant – the “fixer” – was given an adapted task. Throughout the study, fixers were not told to re-create patterns, but to fix patterns to resemble a target grid exactly. Fixers in this experiment viewed the same target grid as learners, but instead of seeing an empty input grid, they saw a grid with 10 elements filled in – the elements that the previous learner had submitted. The participant could then edit the 10 items’ positions. There was no “reset” button during this task, so produced patterns reflect participants’ initial instincts.

In Experiment 3, a generation consisted of a learner, who re-created the target grid, and a fixer, who then received the same target grid as well as the learner’s input grid to edit. The fixer’s edited pattern was used as the target grid for the subsequent generation.

Approximately 8% ($n=71$) of participants in Experiment 3 were excluded from analysis due to failure to meet accuracy requirements on the practice trials or failure to select the necessary number of targets on one or more experimental trials. This resulted in a total of 960 participants included in the analysis.

Analysis and Results

As in Experiments 1 and 2, our primary measures of analysis were accuracy and complexity. These measures were computed using the same methods as in the previous experiments.




Figure 3: In the dyad task, reproduction accuracy stays relatively constant across generations. Fixers have significantly higher accuracies than learners. CHANGE POINT SIZE

Fixers and learners had significantly different pattern reproduction accuracies 3. According to a linear mixed-effects model (need to put in formula? Did an lmer with gen and condition (learner/fixer) as fixed effects, with all the same random effects), the accuracies between groups were significantly different ($\beta_{\text{condition-child}} = -0.076$, $t = -8.956$, $p < .001$). The fixers' transmission accuracies did not increase significantly over generations ($\beta_{\text{fixers}} = -0.011$, $t = -1.075$, $p = .283$), while the accuracy of the learners showed a marginally significant increase ($\beta_{\text{learners}} = 0.019$, $t = 1.855$, $p = .064$).

4 shows the relationship between the complexity of fixers' and learners' patterns. In each generation, the learner decreases the complexity of the pattern, and the fixer is able to compensate for some of this loss. AS in Experiment 2, we fit an exponential model to the data. Both conditions show decreases in pattern complexity over generations ($\beta_{\text{learners}} = -0.032$, $t = -4.739$, $p < .001$; $\beta_{\text{fixers}} = -0.021$, $t = -4.12$, $p < .001$), although the effect of generation is stronger for learners compared to fixers ($\beta_{\text{generation}} = -5.236$, $t = -9.363$, $p < .001$). These results hold true for all three measures of complexity (Do i need to report all of these stats??).

5 shows that the presence of an editor does help retain complexity in the grid patterns. The addition of a fixer into the task allowed a higher degree of complexity to be retained in the language over time ($\beta_{\text{condition-child}} = -3.574$, $t = -5.357$, $p < .001$). Additionally, it appeared that the patterns in the dyad condition asymptoted sooner than in the baseline condition (Stats for this?).

General Discussion

Although Experiments 1-3 used a non-linguistic task, we were able to measure change in a culturally-transmitted, learned symbol system. In Experiments 1 and 2, language simplified rapidly and dramatically, reflecting the strong pressure towards simplification in language learning. These findings replicated those of Kempe et al. (2015): when transmitting an artificial language of grid patterns, complexity in the language was lost.

However, the results of Experiment 3 show that this loss




Figure 4: Fixers reintroduce algorithmic complexity which is lost by learners in the dyad condition.




Figure 5: The presence of a fixer in the dyad condition causes a much greater level of algorithmic complexity to be retained across the evolution of a novel language. CHANGE POSITION IN PAPER

is not permanent, but can be reintroduced in the language by way of a secondary participant who helps bring the language towards a stable level of complexity. When the iterated-learning process begins to resemble the true process of language-learning, where children speak with and are subject to correction by those more competent in the language, a lesser amount of complexity was lost during transmission. Additionally, this stable level of complexity is much higher, and is reached earlier in the transmission chain with the help of a fixing participant. This stability in complexity did not mean that the language stopped changing, but that the descriptiveness and transmissibility pressures were in balance. Fixers in Experiment 3 represented caregivers – they were more accurate at reproducing the language, and could therefore be seen as more fluent speakers of the language, just as adults are of their native languages. The learners, on the other hand, had a more difficult task, which greater strained their working memories, similar to the strain on a child language learner who is inundated with new words each day. The fixer's corrected language was passed to the next learner in the chain, representing a child who, after many years of

being corrected by their own parent, becomes a parent, and, in turn, passes their optimal language to the next generation. Due to the higher accuracy by fixers, and therefore greater knowledge of the language, the fixers were able to compensate for some (not all) of the loss in complexity seen by the learners by editing their patterns.

In Experiment 3, the learner's reproduction accuracies were actually increasing over generations. Despite the stable level of complexity, learners found the language easier to reproduce over evolution. Although a high level of descriptiveness was retained in the language, transmissibility was increasing, without the simplicity pressure weighing in. Perhaps the language was becoming stable and complex, with the symbol-patterns changing to be both descriptive and useful, while being easily transmissible. This reflects the optimal evolutionary response to these two competing pressures.

When a caregiver or teacher prevents their child from growing up to believe that "baba" is the word for both "bottle" and "sheep", they are not only helping their individual child become a competent speaker of the language, but they are also re-introducing complexity, and helping the language system as a whole from simplifying to disuse. Data collection is ongoing with children ages 6-8 at a local science museum in both the Experiment 2 and Experiment 3 tasks, in order to investigate whether the pressures of similarity and complexity affect children similarly to how they affect adults in early language-learning conditions.

We do not learn language as passive listeners, who absorb a proportion of the linguistic input they hear. Therefore, we cannot measure language learning only through measuring input, nor through measuring only linguistic output. Languages are both learned and changed through conversations, with feedback and error correction, to evolve to the needs of the language's users. Therefore, we must study language learning in process, to see how it adapts and evolves with communicative interactions.

All code for these analyses are available at

<https://github.com/mcmeyers/iteratedlearning>

Acknowledgements

This research was funded by a James S. McDonnell Foundation Scholar Award to DY.

References

- Bowerman, M. (1982). U shaped behavioral growth. In (pp. 101–145). Academic Press.
- Brentari, D., & Coppola, M. (2013). What sign language creation teaches us about language. *WIREs Cognitive Science*, 4(201-211).
- Chouinard, M. M., & Clark, E. V. (2003). Adult reformulations of child errors as negative evidence. *Journal of Child Language*, 30(3), 637–669.
- Christiansen, M. H., & Kirby, S. (2003). Language evolution. In M. H. Christiansen & S. Kirby (Eds.), (pp. 1–15). Oxford University Press.
- Ellis, N. C. (2008). The dynamics of second language emergence: Cycles of language use, language change, and language acquisition. *The Modern Language Journal*, 92(ii), 232–249.
- Gauvrit, N., Soler-Toscano, F., & Guida, A. (2017). A preference for some types of complexity comment on "perceived beauty of random texture patterns: A preference for complexity". *Acta Psychologica*, 174, 48–53.
- Hudson Kam, C. L., & Newport, E. L. (2005). Regularizing unpredictable variation: The roles of adult and child learners in language formation and change. *Language Learning and Development*, 1(2), 151–195.
- Kempe, V., Gauvrit, N., & Forsyth, D. (2015). Structure emerges faster during cultural transmission in children than in adults. *Cognition*, 136, 247–254.
- Kirby, S., Dowman, M., & Griffiths, T. L. (2007). Innateness and culture in the evolution of language. *Proceedings of the National Academy of Sciences*, 104(12), 5241–5245.
- Kirby, S., Griffiths, T., & Smith, K. (2014). Iterated learning and the evolution of language. *Current Opinion in Neurobiology*, 28, 108–114.
- Lupyan, G., & Dale, R. (2010). Language structure is partly determined by social structure. *PLoS ONE*, 5(1), 1–10.
- Penner, S. G. (1987). Parental responses to grammatical and ungrammatical child utterances. *Child Development*, 58(2), 376–384.
- Raviv, L., & Arnon, I. (2018). Systematicity, but not compositionality: Examining the emergence of linguistic structure in children and adults using iterated learning. *Cognition*, 181, 160–173.
- Romberg, A. R., & Saffran, J. (2010). Statistical learning and language acquisition. *WIREs Cognitive Science*, 1, 906–914.
- Sandler, W., Meir, I., Padden, C., & Aronoff, M. (2005). The emergence of grammar: Systematic structure in a new language. *Proceedings of the National Academy of Sciences*, 102(7), 2661–2665.
- Senghas, A. (2003). Intergenerational influence and ontogenetic development in the emergence of spatial grammar in Nicaraguan sign language. *Cognitive Development*, 18, 511–531.
- Smith, K., & Wonnacott, E. (2010). Eliminating unpredictable variation through iterated learning. *Cognition*, 116, 444–449.
- structure, E. of combinatorial, & signals. (2014). Verhoef, tessa and kirby, simon and de boer, bart. *Journal of Phonetics*, 43, 57–68.
- Zenil, H., Soler-Toscano, F., Dingle, K., & Louis, A. A. (2014). Correlation of automorphism group size and topological properties with program-size complexity evaluations of graphs and complex networks. *Physica A*, 404, 341–358.